

Feed-Forward Neural Networks, Autoencoders, and Generative Models

1 Introduction

In this assignment, we first use simple supervised networks to classify handwritten digits, and will then apply these techniques to generate novel images of digits using Variational Autoencoders and GANs. Afterwards, we will apply information theoretic measures to assess the success of the generative models.

▼ 2 MNIST Classification

```
1 import torch
2 import torchvision.datasets as datasets
3 import torchvision.transforms as transforms
4 import matplotlib.pyplot as plt
5
6 device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
```

▼ 2.1 Downloading MNIST

```
1 mnist_train = datasets.MNIST(root = 'data', train=True, download=True, trans
2 mnist_test = datasets.MNIST(root = 'data', train=False, download=True,transf
```

Downloading <http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz> to
9920512/? [00:20<00:00, 41882769.81it/s]

Extracting data/MNIST/raw/train-images-idx3-ubyte.gz to data/MNIST/raw
Downloading <http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz> to
32768/? [00:00<00:00, 400862.61it/s]

Extracting data/MNIST/raw/train-labels-idx1-ubyte.gz to data/MNIST/raw
Downloading <http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz> to d
1654784/? [00:19<00:00, 141342.26it/s]

Extracting data/MNIST/raw/t10k-images-idx3-ubyte.gz to data/MNIST/raw
Downloading <http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz> to d
8192/? [00:00<00:00, 28643.81it/s]

Extracting data/MNIST/raw/t10k-labels-idx1-ubyte.gz to data/MNIST/raw
Processing...
Done!
/usr/local/lib/python3.6/dist-packages/torchvision/datasets/mnist.py:480: U
return torch.from_numpy(parsed.astype(m[2], copy=False)).view(*s)



```
1 def evaluate(model, evaluation_set, loss_fn):
2     """
3     Evaluates the given model on the given dataset.
4     Returns the percentage of correct classifications out of total classific
5     """
6     with torch.no_grad(): # this disables backpropogation, which makes the n
7         # TODO: Fill in the rest of the evaluation function.
8         losses = []
9         sum_total = 0
10        for data, targets in evaluation_set:
11            data = data.to(device)
12            targets = targets.to(device)
13            model_input = data.view(-1, 784)
14            out = model(model_input)
15            arg_maxed = torch.argmax(out, dim = 1)
16
17            sum_total += (arg_maxed == targets).float().sum()
18            losses.append(loss_fn(out, targets).item())
19            loss = sum(losses) / len(losses)
20            accuracy = 100 * sum_total / len(evaluation_set.dataset)
21        return accuracy, loss
22
23 def train(model, loss_fn, optimizer, train_loader, test_loader):
```

```

24 """
25 This is a standard training loop, which leaves some parts to be filled i
26 INPUT:
27 :param model: an untrained pytorch model
28 :param loss_fn: e.g. Cross Entropy loss of Mean Squared Error.
29 :param optimizer: the model optimizer, initialized with a learning rate.
30 :param training_set: The training data, in a dataloader for easy iterati
31 :param test_loader: The testing data, in a dataloader for easy iteratio
32 """
33 num_epochs = 100 # obviously, this is too many. I don't know what this a
34 train_loss = []
35 train_acc = []
36 test_loss = []
37 test_acc = []
38 for epoch in range(num_epochs):
39     # loop through each data point in the training set
40     for data, targets in train_loader:
41         data = data.to(device)
42         targets = targets.to(device)
43         optimizer.zero_grad()
44
45         # run the model on the data
46         model_input = data.view(-1, 784)# TODO: Turn the 28 by 28 image
47         out = model(model_input)
48
49         # Calculate the loss
50         loss = loss_fn(out,targets)
51
52         # Find the gradients of our loss via backpropogation
53         loss.backward()
54
55         # Adjust accordingly with the optimizer
56         optimizer.step()
57
58     # Give status reports every 100 epochs
59     if epoch % 100==0:
60         print(f" EPOCH {epoch}. Progress: {epoch/num_epochs*100}%." )
61         tr_acc, tr_loss = evaluate(model, train_loader, loss_fn)
62         te_acc, te_loss = evaluate(model, test_loader, loss_fn)
63         train_loss.append(tr_loss)
64         train_acc.append(tr_acc)
65         test_loss.append(te_loss)
66         test_acc.append(te_acc)
67
68
69         print(f" Train accuracy: {tr_acc}. Test accuracy: {te_acc}") #TC
70
71     return train_loss, train_acc, test_loss, test_acc

```

72

▼ 2.2 Logistic Regression

First, we will explore implementing a simple model that can take in images of handwritten digits and classify them from 0 to 9. We will do this using a multinomial logistic regression model trained with gradient descent.

```
1 from torch.nn.functional import softmax
2 import torch.nn.functional as F
3 from torch import optim, nn
4 import torchvision.transforms as transforms
5 import torchvision.datasets as datasets
6 import numpy as np
7 import torch
```

```

1 class LogisticRegression(nn.Module): # initialize a pytorch neural network m
2     def __init__(self): # initialize the model
3         super(LogisticRegression, self).__init__() # call for the parent cla
4         # you can define variables here that apply to the entire model (e.g.
5         # this model only has two parameters: the weight, and the bias.
6         # here's how you can initialize the weight:
7         # W = nn.Parameter(torch.zeros(shape)) # this creates a model parame
8         # ... torch.zeros is much like numpy.zeros, except optimized for bac
9
10        # create a bias variable here
11        self.W = nn.Parameter(torch.zeros((784, 10)), requires_grad = True)
12        self.b = nn.Parameter(torch.zeros((1, 10)), requires_grad = True)
13
14
15    def forward(self, x):
16        """
17        this is the function that will be executed when we call the logistic
18        INPUT:
19            x, an MNIST image represented as a tensor of shape 784
20        OUTPUT:
21            predictions, a tensor of shape 10. If using CrossEntropyLoss, yc
22        """
23        # put the logic here.
24        predictions = torch.matmul(x, self.W) + self.b
25        predictions = softmax(predictions)
26
27        return predictions

```

```


1 model = LogisticRegression().to(device)
2 # initialize the optimizer, and set the learning rate
3 SGD = torch.optim.SGD(model.parameters(), lr = 5e-1)
4 # initialize the loss function. You don't want to use this one, so change it
5 loss_fn = torch.nn.CrossEntropyLoss()
6 batch_size = 128
7 train_loader = torch.utils.data.DataLoader(mnist_train, batch_size=batch_size)
8 test_loader = torch.utils.data.DataLoader(mnist_test, batch_size=batch_size)
9 train_loss, train_acc, test_loss, test_acc = train(model=model, loss_fn=loss_fn,

```

```

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:25: UserWarning:
EPOCH 0. Progress: 0.0%.
Train accuracy: 87.461669921875. Test accuracy: 88.30999755859375
EPOCH 100. Progress: 66.66666666666666%.
Train accuracy: 93.94667053222656. Test accuracy: 92.94999694824219

```



Question 2.2.1. What percentage classification accuracy did your simple network achieve?

~91%

▼ 2.3 Feed-forward Neural Network

```
1 class FeedForwardNet(nn.Module):
2     """ Simple feed forward network with one hidden layer."""
3     # Here, you should place an exact copy of the code from the LogisticRegr
4     # 1. Add another weight and bias vector to represent the hidden layer
5     # 2. In the forward function, add some type of nonlinearity to the output
6     def __init__(self):
7         super(FeedForwardNet, self).__init__()
8         self.W1 = nn.Parameter(torch.randn((784, 128)), requires_grad = True)
9         self.b1 = nn.Parameter(torch.zeros((1, 128)), requires_grad = True)
10        self.W2 = nn.Parameter(torch.randn((128, 10)), requires_grad = True)
11        self.b2 = nn.Parameter(torch.zeros((1, 10)), requires_grad = True)
12
13    def forward(self, x):
14        predictions = torch.matmul(x, self.W1) + self.b1
15        predictions = F.relu(predictions)
16        predictions = torch.matmul(predictions, self.W2) + self.b2
17        predictions = softmax(predictions)
18        return predictions
```

```

1 model = FeedForwardNet().to(device)
2 # initialize the optimizer, and set the learning rate
3 adam = torch.optim.Adam(model.parameters(), lr = 3e-3) # This is absurdly hi
4 # initialize the loss function. You don't want to use this one, so change it
5 loss_fn = torch.nn.CrossEntropyLoss()
6 batch_size = 128
7 train_loader = torch.utils.data.DataLoader(mnist_train, batch_size=batch_siz
8 test_loader = torch.utils.data.DataLoader(mnist_test, batch_size=batch_size,
9 train_loss, train_acc, test_loss, test_acc = train(model = model, loss_fn = l

```

```


/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:17: UserWarning:
EPOCH 0. Progress: 0.0%.
Train accuracy: 67.92500305175781. Test accuracy: 68.18999481201172
EPOCH 100. Progress: 20.0%.
Train accuracy: 88.91166687011719. Test accuracy: 87.9000015258789
EPOCH 200. Progress: 40.0%.
Train accuracy: 89.20833587646484. Test accuracy: 87.94999694824219
EPOCH 300. Progress: 60.0%.
Train accuracy: 89.44667053222656. Test accuracy: 88.04000091552734
EPOCH 400. Progress: 80.0%.
Train accuracy: 89.66666412353516. Test accuracy: 88.29999542236328

```

```

1 epochs = range(0, len(train_loss))
2 plt.plot(epochs, train_loss, 'g', label='Train loss')
3 plt.plot(epochs, test_loss, 'b', label='Test loss')
4 plt.title('Train and Test loss')
5 plt.xlabel('Epochs')
6 plt.ylabel('Loss')
7 plt.legend()
8 plt.show()

```



```

1 from sklearn.metrics import confusion_matrix as conf_mat
2 labels = []
3 outs = []
4 with torch.no_grad():
5     for i, (inputs, classes) in enumerate(test_loader):
6         inputs = inputs.to(device)
7         classes = classes.to(device)
8         outputs = model(inputs.view(-1, 784))
9         preds = torch.argmax(outputs, 1)
10        labels.extend(classes.detach().cpu().numpy())
11        outs.extend(preds.detach().cpu().numpy())
12
13 conf_mat(labels, outs)

```

```

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:17: UserWarning:
array([[ 973,    0,    1,    1,    0,    0,    1,    2,    2,    0],
       [    0, 1123,    3,    1,    0,    0,    5,    1,    2,    0],
       [    2,    2, 1006,    6,    1,    0,    4,    6,    5,    0],
       [    0,    0,    6, 980,    0,   13,    0,    3,    8,    0],
       [    1,    1,    2,    0, 972,    0,    6,    0,    0,    0],
       [    4,    0,    0,    7,    3, 861,    9,    1,    7,    0],
       [    3,    1,    1,    1,    1,    2, 946,    0,    3,    0],
       [    1,    5,    8,    1,    1,    1,    0, 1006,    5,    0],
       [    5,    1,    1,    4,    1,    0,    2,    5, 955,    0],
       [   13,   11,    3, 108, 595,   69,    0, 105, 105,    0]])

```


Question 2.3.1. What percentage classification accuracy does this more complex network achieve?

89%

Question 2.3.2. Create a plot of the training and test error vs the number of iterations. How many iterations are sufficient to reach good performance?

around 150 epochs seems to be sufficient

Question 2.3.3. Print the confusion matrix showing which digits were misclassified, and what they were misclassified as. What numbers are frequently confused with one another by your model?

For some reason there seems to be a problem with the 9 digit. It was never classifier correctly and it is often misclassified with the 4 digit.

Question 2.3.4. Experiment with the learning rate, optimizer and activation function of your network. Report the best accuracy and briefly describe the training scheme that reached this accuracy.

Using relu activation functions with an Adam optimizer this accuracy was achieved. It worked best with $1e-3$ learning rate. I tried with various learning rates and it seemed Adam with this learning rate worked best

▼ 3 Autoencoder

```

1 class Autoencoder(nn.Module):
2     def __init__(self):
3         super(Autoencoder, self).__init__()
4         self.enc_lin1 = nn.Linear(784, 1000)
5         self.enc_lin2 = nn.Linear(1000, 500)
6         self.enc_lin3 = nn.Linear(500, 250)
7         self.enc_lin4 = nn.Linear(250, 2)
8
9         self.dec_lin1 = nn.Linear(2, 250)
10        self.dec_lin2 = nn.Linear(250, 500)
11        self.dec_lin3 = nn.Linear(500, 1000)
12        self.dec_lin4 = nn.Linear(1000, 784)
13
14
15    def encode(self, x):
16        x = F.tanh(self.enc_lin1(x))
17        x = F.tanh(self.enc_lin2(x))
18        x = F.tanh(self.enc_lin3(x))
19        x = self.enc_lin4(x)
20
21        # ... additional layers, plus possible nonlinearities.
22        return x
23
24    def decode(self, z):
25        # ditto, but in reverse
26        z = F.tanh(self.dec_lin1(z))
27        z = F.tanh(self.dec_lin2(z))
28        z = F.tanh(self.dec_lin3(z))
29        z = F.sigmoid(self.dec_lin4(z))
30
31        return z
32
33    def forward(self, x):
34        z = self.encode(x)
35        return self.decode(z)
36

```



```

1 import matplotlib.pyplot as plt
2
3 def evaluate_ae(model, evaluation_set, loss_fn):
4     """
5     Evaluates the given model on the given dataset.
6     Returns the percentage of correct classifications out of total classific
7     """
8     with torch.no_grad(): # this disables backpropogation, which makes the n
9         # TODO: Fill in the rest of the evaluation function.
10        losses = []

```

```

11     sum_total = 0
12     for ind, (data, targets) in enumerate(evaluation_set):
13         data = data.to(device)
14         model_input = data.view(-1, 784)
15         out = model(model_input)
16         # if ind == 0:
17         #     visualise_output(model_input, model)
18         sum_total += (out == model_input).float().sum()
19         losses.append(loss_fn(out, model_input).item())
20     loss = sum(losses) / len(losses)
21     accuracy = 100 * sum_total / len(evaluation_set.dataset)
22     return accuracy, loss
23
24 def train_ae(model, loss_fn, optimizer, train_loader, test_loader):
25     """
26     This is a standard training loop, which leaves some parts to be filled i
27     INPUT:
28     :param model: an untrained pytorch model
29     :param loss_fn: e.g. Cross Entropy loss or Mean Squared Error.
30     :param optimizer: the model optimizer, initialized with a learning rate.
31     :param training_set: The training data, in a dataloader for easy iterati
32     :param test_loader: The testing data, in a dataloader for easy iterator
33     """
34     num_epochs = 500 # obviously, this is too many. I don't know what this a
35     train_loss = []
36     train_acc = []
37     test_loss = []
38     test_acc = []
39
40     to_graph = {}
41     while True:
42         for data, targets in test_loader:
43             for d, t in zip(data, targets):
44                 if int(t) not in to_graph.keys():
45                     to_graph[int(t)] = d
46                     if len(to_graph) == 10:
47                         break
48             if len(to_graph) == 10:
49                 break
50         if len(to_graph) == 10:
51             break
52
53     for epoch in range(num_epochs):
54         # loop through each data point in the training set
55         for data, targets in train_loader:
56             optimizer.zero_grad()
57             data = data.to(device)
58             # run the model on the data

```

```

59     model_input = data.view(-1, 784) # TODO: Turn the 28 by 28 image
60     out = model(model_input)
61
62     # Calculate the loss
63     loss = loss_fn(out, model_input)
64
65     # Find the gradients of our loss via backpropagation
66     loss.backward()
67
68     # Adjust accordingly with the optimizer
69     optimizer.step()
70
71     # Give status reports every 100 epochs
72     if epoch % 100 == 0:
73         print(f" EPOCH {epoch}. Progress: {epoch/num_epochs*100}%. ")
74         tr_acc, tr_loss = evaluate_ae(model, train_loader, loss_fn)
75         te_acc, te_loss = evaluate_ae(model, test_loader, loss_fn)
76         train_loss.append(tr_loss)
77         train_acc.append(tr_acc)
78         test_loss.append(te_loss)
79         test_acc.append(te_acc)
80
81
82
83         # im = next(iter(train_loader))[0][0]
84         # plt.imshow(im.squeeze())
85         # plt.show()
86         # # import pdb; pdb.set_trace()
87
88         # with torch.no_grad():
89         #     plt.imshow(torch.reshape(model(im.view(-1, 784)), (28,28)).r
90         #     plt.show()
91
92
93         print(f" Train Loss: {tr_loss}. Test Loss: {te_loss}") #TODO: in
94
95     for label, data in to_graph.items():
96         plt.imshow(data.squeeze())
97         plt.show()
98         with torch.no_grad():
99             plt.imshow(torch.reshape(model(data.view(-1, 784)), (28,28)).numpy())
100             plt.show()
101     return train_loss, train_acc, test_loss, test_acc

```

▼ 3.1 MNIST

```

1 model = Autoencoder().to(device)
2 # initialize the optimizer, and set the learning rate
3 optimizer = torch.optim.Adam(model.parameters(), lr = 0.003) # This is absur
4 # initialize the loss function. You don't want to use this one, so change it
5 loss_fn = torch.nn.MSELoss()
6 batch_size = 128
7 train_loader = torch.utils.data.DataLoader(mnist_train, batch_size=batch_siz
8 test_loader = torch.utils.data.DataLoader(mnist_test, batch_size=batch_size,
9 train_ae(model = model, loss_fn = loss_fn, optimizer = optimizer, train_loade

```

```

/usr/local/lib/python3.6/dist-packages/torch/nn/functional.py:1628: UserWar
  warnings.warn("nn.functional.tanh is deprecated. Use torch.tanh instead.")
/usr/local/lib/python3.6/dist-packages/torch/nn/functional.py:1639: UserWar
  warnings.warn("nn.functional.sigmoid is deprecated. Use torch.sigmoid ins
EPOCH 0. Progress: 0.0%.
Train Loss: 0.04769748952120606. Test Loss: 0.047753466222482395
EPOCH 100. Progress: 20.0%.
EPOCH 200. Progress: 40.0%.
Train Loss: 0.046536173179014914. Test Loss: 0.04638722930339318
EPOCH 300. Progress: 60.0%.
Train Loss: 0.051131408375654136. Test Loss: 0.05109376780971696
EPOCH 400. Progress: 80.0%.
Train Loss: 0.0534530708641767. Test Loss: 0.05327377793721006

```

```

-----
NameError                                Traceback (most recent call last)
<ipython-input-8-5c402a189ed2> in <module>()
      7 train_loader = torch.utils.data.DataLoader(mnist_train,
batch_size=batch_size, shuffle=True)
      8 test_loader = torch.utils.data.DataLoader(mnist_test,
batch_size=batch_size, shuffle=True)
----> 9 train_ae(model = model, loss_fn = loss_fn, optimizer = optimizer,
train_loader = train_loader, test_loader = test_loader)

<ipython-input-7-fe3d57a57411> in train_ae(model, loss_fn, optimizer,
train_loader, test_loader)
     92
     93     for label, data in to_graph.items():
--> 94         plt.imshow(data.squeeze())
     95         plt.show()
     96         with torch.no_grad():

```

```


1 import matplotlib.pyplot as plt
2
3 to_graph = {}
4 while True:
5     for data, targets in train_loader:
6         for d, t in zip(data, targets):
7             if int(t) not in to_graph.keys():
8                 to_graph[int(t)] = d

```

```


9     if len(to_graph) == 10:
10         break
11     if len(to_graph) == 10:
12         break
13     if len(to_graph) == 10:
14         break
15
16 for label, data in to_graph.items():
17     plt.imshow(data.squeeze())
18     plt.show()
19     with torch.no_grad():
20         data = data.to(device)
21         plt.imshow(torch.reshape(model(data.view(-1, 784)), (28,28)).cpu().numpy)
22         plt.show()


```



/usr/local/lib/python3.6/dist-packages/torch/nn/functional.py:1628: UserWarning: nn.functional.tanh is deprecated. Use torch.tanh instead.

/usr/local/lib/python3.6/dist-packages/torch/nn/functional.py:1639: UserWarning: nn.functional.sigmoid is deprecated. Use torch.sigmoid instead.






```
1 import seaborn as sns
2 import pandas as pd
3 import matplotlib.pyplot as plt
```

```

1
2
3 thousand_loader = torch.utils.data.DataLoader(mnist_train, batch_size = 1000
4 for data, target in thousand_loader:
5     with torch.no_grad():
6         data = data.to(device)
7         model_input = data.view(-1, 784)# TODO: Turn the 28 by 28 image tensors
8         out = model.encode(model_input).detach().cpu().numpy()
9         labels = target
10        break
11
12 df = {'x': out[:, 0], 'y': out[:, 1], 'label': labels.numpy()}
13
14
15
16 data = pd.DataFrame(df)
17 facet = sns.lmplot(data=data, x='x', y='y', hue='label',
18                    fit_reg=False, legend=False)
19
20 #add a legend
21 leg = facet.ax.legend(bbox_to_anchor=[1, 0.75],
22                       title="label", fancybox=True)
23

```

/usr/local/lib/python3.6/dist-packages/torch/nn/functional.py:1628: UserWarning: warnings.warn("nn.functional.tanh is deprecated. Use torch.tanh instead.")



Question 3.1.1. Do the colors easily separate, or are they all clumped together? Which numbers are frequently embedded close together, and what does this mean?

The colors easily separate, but it does not seem like there are any distinct clusters. Perhaps with optimized training there would be. The 1s and 0s seem to be pretty distinct, but something like an 8 which has many lines, seems to be scattered throughout. It also seems like 2s and 6s are somewhat close, which means the network learned these embeddings to have similar values.

Question 3.1.2. How realistic were the images you generated by interpolating between points in the latent space? Can you think of a better way to generate images with an autoencoder?

The generated images were not too realistic. Perhaps with more dimensions between the encoder and decoder, we could see better results.

▼ 3.2 Biological Data: Retinal Bipolar Dataset

```
1 import pickle
2 with open('retinal-bipolar-data.pickle', 'rb') as f:
3     retinal = pickle.load(f)
4 with open('retinal-bipolar-metadata.pickle', 'rb') as f:
5     retinal_meta = pickle.load(f)

1 retinal.shape

(21552, 15524)

1 msk = np.random.rand(len(retinal)) < 0.8
2
3 ret_train = torch.Tensor(retinal[msk].values.astype(np.float32))
4 ret_tr_y = torch.Tensor(retinal_meta[msk]['CLUSTER'].values.astype(np.float32))
5 ret_test = torch.Tensor(retinal[~msk].values.astype(np.float32))
6 ret_te_y = torch.Tensor(retinal_meta[~msk]['CLUSTER'].values.astype(np.float32))

1 ret_train = torch.utils.data.TensorDataset(ret_train, ret_tr_y)
2 ret_test = torch.utils.data.TensorDataset(ret_test, ret_te_y)
```

```

1 class Ret_Autoencoder(Autoencoder):
2     def __init__(self, input_size):
3         super(Autoencoder, self).__init__()
4         self.enc_lin1 = nn.Linear(input_size, 1000)
5         self.enc_lin2 = nn.Linear(1000, 500)
6         self.enc_lin3 = nn.Linear(500, 250)
7         self.enc_lin4 = nn.Linear(250, 2)
8
9         self.dec_lin1 = nn.Linear(2, 250)
10        self.dec_lin2 = nn.Linear(250, 500)
11        self.dec_lin3 = nn.Linear(500, 1000)
12        self.dec_lin4 = nn.Linear(1000, input_size)
13
14    def decode(self, z):
15        # ditto, but in reverse
16        z = F.tanh(self.dec_lin1(z))
17        z = F.tanh(self.dec_lin2(z))
18        z = F.tanh(self.dec_lin3(z))
19        z = self.dec_lin4(z)
20
21        return z
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100

```

```

1 import matplotlib.pyplot as plt
2
3 def evaluate_ret_ae(model, evaluation_set, loss_fn):
4     """
5     Evaluates the given model on the given dataset.
6     Returns the percentage of correct classifications out of total classifications.
7     """
8     with torch.no_grad(): # this disables backpropagation, which makes the n
9         # TODO: Fill in the rest of the evaluation function.
10        losses = []
11        sum_total = 0
12        for ind, (data, targets) in enumerate(evaluation_set):
13            data = data.to(device)
14            model_input = data
15            out = model(model_input)
16            # if ind == 0:
17            #     visualise_output(model_input, model)
18            sum_total += (out == model_input).float().sum()
19            losses.append(loss_fn(out, model_input).item())
20        loss = sum(losses) / len(losses)
21        accuracy = 100 * sum_total / len(evaluation_set.dataset)
22        return accuracy, loss
23
24 def train_ret_ae(model, loss_fn, optimizer, train_loader, test_loader):

```

```


25  """
26  This is a standard training loop, which leaves some parts to be filled i
27  INPUT:
28  :param model: an untrained pytorch model
29  :param loss_fn: e.g. Cross Entropy loss of Mean Squared Error.
30  :param optimizer: the model optimizer, initialized with a learning rate.
31  :param training_set: The training data, in a dataloader for easy iterati
32  :param test_loader: The testing data, in a dataloader for easy iterator
33  """
34  num_epochs = 500 # obviously, this is too many. I don't know what this a
35  train_loss = []
36  train_acc = []
37  test_loss = []
38  test_acc = []
39
40  for epoch in range(num_epochs):
41      # loop through each data point in the training set
42      for data, targets in train_loader:
43          optimizer.zero_grad()
44          data = data.to(device)
45          # run the model on the data
46          model_input = data
47          out = model(model_input)
48
49          # Calculate the loss
50          loss = loss_fn(out, model_input)
51
52          # Find the gradients of our loss via backpropogation
53          loss.backward()
54
55          # Adjust accordingly with the optimizer
56          optimizer.step()
57
58      # Give status reports every 100 epochs
59      if epoch % 100==0:
60          print(f" EPOCH {epoch}. Progress: {epoch/num_epochs*100}%.")
61          tr_acc, tr_loss = evaluate_ret_ae(model, train_loader, loss_fn)
62          te_acc, te_loss = evaluate_ret_ae(model, test_loader, loss_fn)
63          train_loss.append(tr_loss)
64          train_acc.append(tr_acc)
65          test_loss.append(te_loss)
66          test_acc.append(te_acc)
67
68          print(f" Train Loss: {tr_loss}. Test Loss: {te_loss}") #TODO: in
69
70  return train_loss, train_acc, test_loss, test_acc

```

```
1 model = Ret_Autoencoder(input_size = 15524).to(device)
2 # initialize the optimizer, and set the learning rate
3 optimizer = torch.optim.Adam(model.parameters(), lr = 0.003) # This is absur
4 # initialize the loss function. You don't want to use this one, so change it
5 loss_fn = torch.nn.MSELoss()
6 batch_size = 128
7 train_loader = torch.utils.data.DataLoader(ret_train, batch_size=batch_size,
8 test_loader = torch.utils.data.DataLoader(ret_test, batch_size=batch_size, s
9 train_ret_ae(model = model, loss_fn = loss_fn, optimizer = optimizer, train_l

1 thousand_loader = torch.utils.data.DataLoader(ret_test, batch_size = 1000, s
2 for data, target in thousand_loader:
3     with torch.no_grad():
4         data = data.to(device)
5         model_input = data
6         out = model.encode(model_input).detach().cpu().numpy()
7         labels = target
8         break
9
10 df = {'x': out[:, 0], 'y': out[:, 1], 'label': labels.numpy()}
11
12
13
14 data = pd.DataFrame(df)
15 facet = sns.lmplot(data=data, x='x', y='y', hue='label',
16                     fit_reg=False, legend=False)
17
18 #add a legend
19 leg = facet.ax.legend(bbox_to_anchor=[1, 0.75],
20                        title="label", fancybox=True)
```

```
/usr/local/lib/python3.6/dist-packages/torch/nn/functional.py:1628: UserWarning:
  warnings.warn("nn.functional.tanh is deprecated. Use torch.tanh instead.")
```



Question 3.2.1. How many clusters are visible in the embedding? Do they correspond to the cluster labels?

It seems like there was only three clusters. The y dimension seems to have very little variance. It does seem, however, that there is a trend, where numbers found the same x value and ended up in three different bins.

▼ 4 Generative Models

▼ 4.1 The Variational Autoencoder

```
1 ##### vae.py
2
3 import torch
4 import torch.utils.data
5 from torch import nn, optim
6 from torch.nn import functional as F
7 from torchvision import datasets, transforms
8 from torchvision.utils import save_image
9
10 train_loader = torch.utils.data.DataLoader(mnist_train, batch_size=batch_size)
11 test_loader = torch.utils.data.DataLoader(mnist_test, batch_size=batch_size)
12
13 batch_size = 128
14 epochs = 10
15 seed = 1
16 log_interval = 10
17
18
19 class VAE(nn.Module):
20     def __init__(self):
21         super(VAE, self).__init__()
22
23         self.fc1 = nn.Linear(784, 400)
24         self.fc21 = nn.Linear(400, 20)
25         self.fc22 = nn.Linear(400, 20)
26         self.fc3 = nn.Linear(20, 400)
27         self.fc4 = nn.Linear(400, 784)
28
29     def encode(self, x):
30         h1 = F.relu(self.fc1(x))
31         return self.fc21(h1), self.fc22(h1)
32
33     def reparameterize(self, mu, logvar):
34         std = torch.exp(0.5*logvar)
35         eps = torch.randn_like(std)
36         return mu + eps*std
37
38     def decode(self, z):
39         h3 = F.relu(self.fc3(z))
40         return torch.sigmoid(self.fc4(h3))
41
42     def forward(self, x):
43         mu, logvar = self.encode(x.view(-1, 784))
44         z = self.reparameterize(mu, logvar)
45         return self.decode(z), mu, logvar
46
47
48
```

```

49
50
51 def VAE_loss_function(recon_x, x, mu, logvar):
52     # TO DO: Implement reconstruction + KL divergence losses summed over all
53
54     # see lecture 12 slides for more information on the VAE loss function
55     # for additional information on computing KL divergence
56     # see Appendix B from VAE paper:
57     # Kingma and Welling. Auto-Encoding Variational Bayes. ICLR, 2014
58     # https://arxiv.org/abs/1312.6114
59
60     # x = x.squeeze(1)
61     recon_loss = F.binary_cross_entropy(recon_x, x.view(-1, 784), reduction=
62     KLD = -0.5 * torch.sum(1 + logvar - mu.pow(2) - logvar.exp())
63     return recon_loss + KLD
64
65
66 def train(epoch):
67     model.train()
68     train_loss = 0
69     for batch_idx, (data, _) in enumerate(train_loader):
70         data = data.to(device)
71         optimizer.zero_grad()
72         recon_batch, mu, logvar = model(data)
73         loss = VAE_loss_function(recon_batch, data, mu, logvar)
74         loss.backward()
75         train_loss += loss.item()
76         optimizer.step()
77         if batch_idx % log_interval == 0:
78             print('Train Epoch: {} [{}/{} ({:.0f}%)]\tLoss: {:.6f}'.format(
79                 epoch, batch_idx * len(data), len(train_loader.dataset),
80                 100. * batch_idx / len(train_loader),
81                 loss.item() / len(data)))
82
83     print('====> Epoch: {} Average loss: {:.4f}'.format(
84         epoch, train_loss / len(train_loader.dataset)))
85
86
87 def test(epoch):
88     model.eval()
89     test_loss = 0
90     with torch.no_grad():
91         for i, (data, _) in enumerate(test_loader):
92             data = data.to(device)
93             recon_batch, mu, logvar = model(data)
94             test_loss += VAE_loss_function(recon_batch, data, mu, logvar).it
95             if i == 0:
96                 n = min(data.size(0), 8)

```

```

98         comparison = torch.cat([data_batch.view(batch_size, 1, 28, 28)
99         save_image(comparison.cpu(),
100             'results/reconstruction_' + str(epoch) + '.png', nr
101
102     test_loss /= len(test_loader.dataset)
103     print('====> Test set loss: {:.4f}'.format(test_loss))
104
105
106

```

```
1 !mkdir results
```

```

1 #if name == main
2 train_loader = torch.utils.data.DataLoader(mnist_train, batch_size=batch_size)
3 test_loader = torch.utils.data.DataLoader(mnist_test, batch_size=batch_size,
4 device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
5 model = VAE().to(device)
6 optimizer = optim.Adam(model.parameters(), lr=1e-3)
7 for epoch in range(1, epochs + 1):
8     train(epoch)
9     test(epoch)
10     with torch.no_grad():
11         sample = torch.randn(64, 20).to(device)
12         sample = model.decode(sample).cpu()
13         save_image(sample.view(64, 1, 28, 28),
14             'results/sample_' + str(epoch) + '.png')

```

```

Train Epoch: 1 [0/60000 (0%)]    Loss: 547.250000
Train Epoch: 1 [1280/60000 (2%)]    Loss: 304.746094
Train Epoch: 1 [2560/60000 (4%)]    Loss: 233.603378
Train Epoch: 1 [3840/60000 (6%)]    Loss: 221.929214
Train Epoch: 1 [5120/60000 (9%)]    Loss: 218.519470
Train Epoch: 1 [6400/60000 (11%)]    Loss: 215.135056
Train Epoch: 1 [7680/60000 (13%)]    Loss: 207.176376
Train Epoch: 1 [8960/60000 (15%)]    Loss: 203.229523
Train Epoch: 1 [10240/60000 (17%)]    Loss: 188.281433
Train Epoch: 1 [11520/60000 (19%)]    Loss: 188.851898
Train Epoch: 1 [12800/60000 (21%)]    Loss: 182.265137
Train Epoch: 1 [14080/60000 (23%)]    Loss: 179.018768
Train Epoch: 1 [15360/60000 (26%)]    Loss: 177.512329
Train Epoch: 1 [16640/60000 (28%)]    Loss: 162.004517
Train Epoch: 1 [17920/60000 (30%)]    Loss: 162.069183
Train Epoch: 1 [19200/60000 (32%)]    Loss: 163.561600
Train Epoch: 1 [20480/60000 (34%)]    Loss: 155.700745
Train Epoch: 1 [21760/60000 (36%)]    Loss: 158.180450
Train Epoch: 1 [23040/60000 (38%)]    Loss: 162.674438
Train Epoch: 1 [24320/60000 (41%)]    Loss: 154.532822
Train Epoch: 1 [25600/60000 (43%)]    Loss: 153.583633
Train Epoch: 1 [26880/60000 (45%)]    Loss: 154.731003

```



```

Train Epoch: 1 [28160/60000 (47%)] Loss: 149.200836
Train Epoch: 1 [29440/60000 (49%)] Loss: 150.456726
Train Epoch: 1 [30720/60000 (51%)] Loss: 146.824585
Train Epoch: 1 [32000/60000 (53%)] Loss: 143.879807
Train Epoch: 1 [33280/60000 (55%)] Loss: 152.822159
Train Epoch: 1 [34560/60000 (58%)] Loss: 143.549591
Train Epoch: 1 [35840/60000 (60%)] Loss: 141.609024
Train Epoch: 1 [37120/60000 (62%)] Loss: 142.283875
Train Epoch: 1 [38400/60000 (64%)] Loss: 140.511612
Train Epoch: 1 [39680/60000 (66%)] Loss: 142.047607
Train Epoch: 1 [40960/60000 (68%)] Loss: 139.073242
Train Epoch: 1 [42240/60000 (70%)] Loss: 140.667725
Train Epoch: 1 [43520/60000 (72%)] Loss: 142.093582
Train Epoch: 1 [44800/60000 (75%)] Loss: 131.476868
Train Epoch: 1 [46080/60000 (77%)] Loss: 137.830429
Train Epoch: 1 [47360/60000 (79%)] Loss: 141.102036
Train Epoch: 1 [48640/60000 (81%)] Loss: 128.376373
Train Epoch: 1 [49920/60000 (83%)] Loss: 135.808212
Train Epoch: 1 [51200/60000 (85%)] Loss: 133.896515
Train Epoch: 1 [52480/60000 (87%)] Loss: 133.259933
Train Epoch: 1 [53760/60000 (90%)] Loss: 127.026588
Train Epoch: 1 [55040/60000 (92%)] Loss: 129.372818
Train Epoch: 1 [56320/60000 (94%)] Loss: 132.166794
Train Epoch: 1 [57600/60000 (96%)] Loss: 127.940308
Train Epoch: 1 [58880/60000 (98%)] Loss: 125.580383
====> Epoch: 1 Average loss: 164.6744
====> Test set loss: 128.0015
Train Epoch: 2 [0/60000 (0%)] Loss: 131.774231
Train Epoch: 2 [1280/60000 (2%)] Loss: 129.155136
Train Epoch: 2 [2560/60000 (4%)] Loss: 123.548080
Train Epoch: 2 [3840/60000 (6%)] Loss: 126.043655
Train Epoch: 2 [5120/60000 (9%)] Loss: 125.900391
Train Epoch: 2 [6400/60000 (11%)] Loss: 125.306747
Train Epoch: 2 [7680/60000 (13%)] Loss: 124.910072
Train Epoch: 2 [8960/60000 (15%)] Loss: 124.997482
Train Epoch: 2 [10240/60000 (17%)] Loss: 123.886559
Train Epoch: 2 [11520/60000 (19%)] Loss: 122.602714

```

1

```

-----
IndexError                                Traceback (most recent call last)
<ipython-input-26-352892a87196> in <module>()
----> 1 to_graph[1][1]

```

```

IndexError: index 1 is out of bounds for dimension 0 with size 1

```



SEARCH STACK OVERFLOW


```

1 to_graph = {}
2 while True:
3   for data, targets in test_loader:

```

```
4     for d, t in zip(data, targets):
5         if int(t) not in to_graph.keys():
6             to_graph[int(t)] = d[0]
7         if len(to_graph) == 10:
8             break
9     if len(to_graph) == 10:
10        break
11 if len(to_graph) == 10:
12     break
13
14 for label, data in to_graph.items():
15     plt.imshow(data.squeeze())
16     plt.show()
17     with torch.no_grad():
18         data = data.to(device)
19         z, _, _ = model(data.view(-1, 784))
20         plt.imshow(torch.reshape(z, (28,28)).cpu().numpy())
21         plt.show()
```





```
1 !pip install phate scprep
```

```
Requirement already satisfied: phate in /usr/local/lib/python3.6/dist-packa
Requirement already satisfied: scprep in /usr/local/lib/python3.6/dist-pack
Requirement already satisfied: tasklogger>=1.0 in /usr/local/lib/python3.6/
Requirement already satisfied: numpy>=1.16.0 in /usr/local/lib/python3.6/di
Requirement already satisfied: scipy>=1.1.0 in /usr/local/lib/python3.6/dis
Requirement already satisfied: Deprecated in /usr/local/lib/python3.6/dist-
Requirement already satisfied: s-gd2>=1.5 in /usr/local/lib/python3.6/dist-
Requirement already satisfied: future in /usr/local/lib/python3.6/dist-pack
Requirement already satisfied: scikit-learn>=0.20.0 in /usr/local/lib/pytho
Requirement already satisfied: graphtools>=1.3.1 in /usr/local/lib/python3.
Requirement already satisfied: matplotlib>=3.0 in /usr/local/lib/python3.6/
Requirement already satisfied: decorator>=4.3.0 in /usr/local/lib/python3.6
Requirement already satisfied: pandas>=0.25 in /usr/local/lib/python3.6/dis
Requirement already satisfied: wrapt<2,>=1.10 in /usr/local/lib/python3.6/d
Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.6/dis
Requirement already satisfied: pygsp>=0.5.1 in /usr/local/lib/python3.6/dis
Requirement already satisfied: cycycler>=0.10 in /usr/local/lib/python3.6/dis
Requirement already satisfied: python-dateutil>=2.1 in /usr/local/lib/pytho
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in
Requirement already satisfied: pytz>=2017.2 in /usr/local/lib/python3.6/dis
Requirement already satisfied: six in /usr/local/lib/python3.6/dist-package
```

```
1 import phate
2 import scprep
3 thousand_loader = torch.utils.data.DataLoader(mnist_train, batch_size = 1000)
4 for data, target in thousand_loader:
5     with torch.no_grad():
6         data = data.to(device)
7         model_input = data.view(-1, 784)# TODO: Turn the 28 by 28 image tensors
8         mu, log = model.encode(model_input)
9         out = model.reparameterize(mu, log)
10        out = out.detach().cpu().numpy()
11        labels = target
12        break
13
14 subsample_data_pc, subsample_meta = scprep.select.subsample(out, labels, n =
15 data_phate = phate.PHATE().fit_transform(subsample_data_pc)
16
17 scprep.plot.scatter2d(data_phate, c=list(subsample_meta), figsize=(15,5), le
18
19 # df = {'x': out[:, 0], 'y': out[:, 1], 'label': labels.numpy()}
20
21
22
23 # data = pd.DataFrame(df)
24 # facet = sns.lmplot(data=data, x='x', y='y', hue='label',
25 #                     fit_reg=False, legend=False)
26
27 # #add a legend
28 # leg = facet.ax.legend(bbox_to_anchor=[1, 0.75],
29 #                       title="label", fancybox=True)
30
```

Calculating PHATE...

Running PHATE on 1000 observations and 20 variables.

Calculating graph and diffusion operator...

Calculating KNN search...

Calculated KNN search in 0.04 seconds.

Calculating affinities...

Calculated affinities in 0.01 seconds.

Calculated graph and diffusion operator in 0.06 seconds.

Calculating optimal t...

Automatically selected $t = 30$

Calculated optimal t in 0.71 seconds.

Calculating diffusion potential...


Calculated diffusion potential in 0.43 seconds.

Calculating metric MDS...

Calculated metric MDS in 1.02 seconds.

Calculated PHATE in 2.23 seconds.

<matplotlib.axes._subplots.AxesSubplot at 0x7f2e58205a90>



Question 4.1.1. How does the VAE's latent space compare to the latent space of your previous autoencoder? Do the generated images have more clarity? Is this most noticeable between or within classes?

The latent space seems to be much better defined for the VAE, and the generated images absolutely have more clarity. It is most noticeable between classes.

Question 4.1.2. In what situations would a VAE be more useful than a vanilla autoencoder, and when would you prefer a vanilla autoencoder to a VAE?

VAEs should perform better on larger datasets, and they perform better to distinguish between classes. I think vanilla AE should be used first, if they are sufficient, then use it, as it is a simpler model. If the data has a lot of similarities, then an AE should be good, but if it is important to make distinctions between categories, then VAE.

Question 4.1.3. The distance between embeddings in your first autoencoder provided some measure of the similarity between digits. To what extent is this preserved, or improved, by the VAE?

This is preserved less. This is by design of the VAE.

▼ 4.2 GANs

```
1 import torch
2 import torch.nn as nn
3 import torchvision.transforms as transforms
4 import torch.optim as optim
5 import torchvision.datasets as datasets
6 from torch.autograd import Variable
7 import imageio
8 import numpy as np
9 import matplotlib
10 import math
11 from torchvision.utils import make_grid, save_image
12 from torch.utils.data import DataLoader
13 from matplotlib import pyplot as plt
14 from tqdm import tqdm
15 matplotlib.style.use('ggplot')
```

```
1 !mkdir outputs
```

```
mkdir: cannot create directory 'outputs': File exists
```

```
1 ##### GAN.py
2 criterion = nn.BCELoss()
3 device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
4 class Generator(nn.Module):
5     def __init__(self, nz):
6         super(Generator, self).__init__()
7         self.nz = nz # the dimension of the random noise used to seed the Ge
8         self.main = nn.Sequential( # nn.sequential is a handy way of combini
9             nn.Linear(self.nz, 256),
10            nn.LeakyReLU(0.2),
11            nn.Linear(256, 512),
12            nn.LeakyReLU(0.2),
13            nn.Linear(512, 1024),
14            nn.LeakyReLU(0.2),
15            nn.Linear(1024, 784),
16            nn.Tanh(),
17        )
18    def forward(self, x):
19        return self.main(x).view(-1, 1, 28, 28)
20
21 class Discriminator(nn.Module):
22     def __init__(self):
23         super(Discriminator, self).__init__()
24         self.n_input = 784
25         self.main = nn.Sequential(
26             nn.Linear(self.n_input, 1024),
27             nn.LeakyReLU(0.2),
28             nn.Dropout(0.3),
29             nn.Linear(1024, 512),
30             nn.LeakyReLU(0.2),
31             nn.Dropout(0.3),
32             nn.Linear(512, 256),
33             nn.LeakyReLU(0.2),
34             nn.Dropout(0.3),
35             nn.Linear(256, 1),
36             nn.Sigmoid(),
37         )
38    def forward(self, x):
39        x = x.view(-1, 784)
40        return self.main(x)
41
42
43 def train_discriminator(optimizer, real_data, fake_data):
```

```

44     """
45     Train the discriminator on a minibatch of data.
46     INPUTS
47         :param optimizer: the optimizer used for training
48         :param real_data: the batch of training data
49         :param fake_data: the data generated by the generator from random noise
50     The discriminator will incur two losses: one from trying to classify the
51     TODO: Fill in this function.
52     It should
53     1. Run the discriminator on the real_data and the fake_data
54     2. Compute and sum the respective loss terms (described in the assignment)
55     3. Backpropagate the loss (e.g. loss.backward()), and perform optimization
56     """
57     # optimizer.zero_grad()
58     fake_out = discriminator(fake_data)
59     real_out = discriminator(real_data)
60     y_fake = Variable(torch.zeros(batch_size, 1, device = device))
61     y_real = Variable(torch.ones(batch_size, 1, device = device))
62     D_real_loss = criterion(real_out, y_real)
63     D_fake_loss = criterion(fake_out, y_fake)
64     # loss = -0.5 * y_real * torch.clamp(torch.log((real_out)), 1e-3, 1e3) -
65     # loss = loss.sum()
66     # import pdb; pdb.set_trace()
67     loss = (D_fake_loss + D_real_loss) / 2
68
69     loss.backward()
70     # optimizer.step()
71
72     # we'll return the loss for book-keeping purposes. (E.g. if you want to
73
74     return loss.item(), optimizer
75
76 def train_generator(optimizer, fake_data):
77     """
78     Performs a single training step on the generator.
79     :param optimizer: the optimizer
80     :param fake_data: forgeries, created by the generator from random noise.
81     :return: the generator's loss
82     TODO: Fill in this function
83     It should
84     1. Run the discriminator on the fake_data
85     2. compute the resultant loss for the generator (as described in the assignment)
86     3. Backpropagate the loss, and perform optimization
87     """
88     # import pdb; pdb.set_trace()
89
90     # with torch.no_grad():
91     # optimizer.zero_grad()

```



```

92     fake_out = discriminator(fake_data)
93     y = Variable(torch.ones(batch_size, 1).to(device))
94     loss = criterion(fake_out, y)
95     # loss = -0.5 * y * torch.clamp(torch.log(fake_out), 1e-3, 1e3)
96     # loss = loss.sum()
97     # import pdb; pdb.set_trace()
98     loss.backward()
99     # optimizer.step()
100     return loss.item(), optimizer
101
102 # import data
103 batch_size = 100
104 train_data = datasets.MNIST(
105     root='../data',
106     train=True,
107     download=True,
108     transform=transforms.ToTensor()
109 )
110 train_loader = DataLoader(train_data, batch_size=batch_size, shuffle=True)
111
112 num_epochs = 1000
113 nz = 25# dimension of random noise
114 generator = Generator(nz)
115 generator = generator.to(device)
116 discriminator = Discriminator()
117 discriminator = discriminator.to(device)
118 g_optimizer = optim.Adam(generator.parameters(), lr=3e-4)
119 d_optimizer = optim.Adam(discriminator.parameters(), lr=3e-4)
120 #TODO: Build a training loop for the GAN
121 # For each epoch, you'll
122 # 1. Loop through the training data. For each batch, feed random noise into
123 # 2. Feed the fake data and real data into the train_discriminator and train
124 # At the end of each epoch, use the below functions to save a grid of genera
125 for epoch in range(num_epochs):
126     for data, _ in train_loader:
127         # perform training
128         data = data.to(device)
129         g_optimizer.zero_grad()
130
131
132         noise = torch.randn((batch_size, nz), device=device)
133         noise = torch.clamp(noise, 1e-8, 1)
134         fake_data = generator(noise)
135         gen_loss, g_optimizer = train_generator(optimizer = g_optimizer, fak
136         g_optimizer.step()
137
138         d_optimizer.zero_grad()
139         noise = torch.randn((batch_size, nz), device=device)

```

```


140         fake_data = torch.randn(1, 1, 1, 1)
141         dis_loss, d_optimizer = train_discriminator(optimizer = d_optimizer,
142
143
144
145         d_optimizer.step()
146     # reshape the image tensors into a grid
147     if epoch % 100 == 0:
148         print('Epoch: {} \t Gen loss: {} \t Dis loss: {}'.format(epoch, gen_loss, dis_loss))
149         print('youre past {}'.format(epoch))
150     generated_img = make_grid(fake_data)
151     # save the generated torch tensor images
152     save_image(generated_img, f"outputs/gen_img{epoch}.png")

```

```

Epoch: 0          Gen loss: 9.10543155670166      Dis loss: 0.48131835460662
youre past 0
Epoch: 100        Gen loss: 1.3410712480545044    Dis loss: 0.46736103296279
youre past 100

```



Question 4.2.1. Which generates more realistic images: your GAN, or your VAE? Why do you think this is?

From what I saw, the VAE generated more realistic images. I assume this is because there is an actual input to the VAE to reconstruct it whereas the GAN generalizes over the noise

Question 4.2.2. Does your GAN appear to generate all digits in equal number, or has it specialized in a smaller number of digits? If so, why might this be?

It seems to generate 0s more than anything. This is probably because so many numbers have curves, 2, 3, 5, 6, 8, 9, 0 that having a number that is essentially the average would be the easiest for the generator to learn. It also seems to make a lot of 3s, possibly because the left side of many numbers is open. interestingly enough, the generator doesn't appear to make many 8s.

```

1 model = LogisticRegression().to(device)
2 # initialize the optimizer, and set the learning rate
3 SGD = torch.optim.SGD(model.parameters(), lr = 5e-2) # This is absurdly high
4 # initialize the loss function. You don't want to use this one, so change it
5 loss_fn = torch.nn.CrossEntropyLoss()
6 batch_size = 128
7 train_loader = torch.utils.data.DataLoader(mnist_train, batch_size=batch_size)
8 test_loader = torch.utils.data.DataLoader(mnist_test, batch_size=batch_size)
9 train_loss, train_acc, test_loss, test_acc = train(model=model, loss_fn=loss_fn)

```


```
1 !unzip to_classify.zip
```

```
Archive:  to_classify.zip
  creating:  to_classify/
  inflating:  to_classify/gen_img906.png
  inflating:  to_classify/gen_img908.png
  inflating:  to_classify/gen_img913.png
  inflating:  to_classify/gen_img914.png
  inflating:  to_classify/gen_img923.png
  inflating:  to_classify/gen_img925.png
  inflating:  to_classify/gen_img928.png
  inflating:  to_classify/gen_img936.png
  inflating:  to_classify/gen_img940.png
  inflating:  to_classify/gen_img943.png
  inflating:  to_classify/gen_img946.png
  inflating:  to_classify/gen_img947.png
  inflating:  to_classify/gen_img951.png
  inflating:  to_classify/gen_img954.png
  inflating:  to_classify/gen_img955.png
  inflating:  to_classify/gen_img956.png
  inflating:  to_classify/gen_img957.png
  inflating:  to_classify/gen_img963.png
  inflating:  to_classify/gen_img968.png
  inflating:  to_classify/gen_img971.png
  inflating:  to_classify/gen_img972.png
  inflating:  to_classify/gen_img973.png
  inflating:  to_classify/gen_img978.png
  inflating:  to_classify/gen_img981.png
  inflating:  to_classify/gen_img999.png
```


```
1 import matplotlib.pyplot as plt
2
3 from PIL import Image
4 import torchvision.transforms.functional as TF
5 import os
6 from torchvision.transforms import Grayscale
7 scaler = Grayscale(1)
8 for filename in os.listdir('to_classify'):
9     print('Generated')
10    image = Image.open(os.path.join('to_classify',filename))
11    plt.imshow(image)
12    plt.show()
13    crop_rectangle = (2, 2, 30, 30) #crop to 28 x 28
14    cropped_im = image.crop(crop_rectangle)
15    cropped_im = scaler(cropped_im) #convert to grayscale
16    print('Cropped/greyscaled')
17    plt.imshow(cropped_im)
18    plt.show()
19    x = TF.to_tensor(cropped_im).to(device)
20    model_input = x.view(-1, 784)
```

```
21 out = model(model_input)
22 arg_maxed = torch.argmax(out, dim = 1)
23 print('predicted: {}'.format(arg_maxed.item()))
24
```

Generated




Cropped/greyscaled

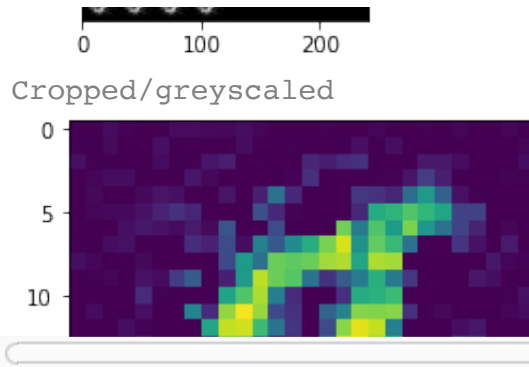


predicted: 0

Generated

```
/usr/local/lib/python3.6/dist-packages/torch/nn/functional.py:1639: UserWarning: nn.functional.sigmoid is deprecated. Use torch.sigmoid instead.
warnings.warn("nn.functional.sigmoid is deprecated. Use torch.sigmoid instead.")
```





▼ 5 Information Theory

```
1 !pip install pyemd
```

```
Requirement already satisfied: pyemd in /usr/local/lib/python3.6/dist-packa
Requirement already satisfied: numpy<2.0.0,>=1.9.0 in /usr/local/lib/python
```

```
1 from time import perf_counter
2 import torch
3 from pyemd import emd
4 from scipy.stats import entropy
5 import numpy as np
6 import pickle
7 import sklearn.neighbors
8 from scipy.spatial import distance_matrix
9 import torchvision.datasets as datasets
10 import torchvision.transforms as transforms
11
12 mnist_train = datasets.MNIST(root = 'data', train=True, download=True, trans
13 mnist_test = datasets.MNIST(root = 'data', train=False, download=True,transf
```

Downloading <http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz> to
 76% 7553024/9912422 [00:00<00:13, 171660.09it/s]

Extracting data/MNIST/raw/train-images-idx3-ubyte.gz to data/MNIST/raw

Downloading <http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz> to
 32768/? [00:00<00:00, 350126.37it/s]

Extracting data/MNIST/raw/train-labels-idx1-ubyte.gz to data/MNIST/raw

Downloading <http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz> to d
 1% 16384/1648877 [00:00<00:12, 126950.32it/s]

Extracting data/MNIST/raw/t10k-images-idx3-ubyte.gz to data/MNIST/raw

Downloading <http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz> to d
 0/? [00:00<?, ?it/s]

Extracting data/MNIST/raw/t10k-labels-idx1-ubyte.gz to data/MNIST/raw

Processing...

Done!

/usr/local/lib/python3.6/dist-packages/torchvision/datasets/mnist.py:480: U
 return torch.from_numpy(parsed.astype(m[2], copy=False)).view(*s)

```

1 def mmd(X,Y, kernel_fn):
2     """
3     Implementation of Maximum Mean Discrepancy.
4     :param X: An n x 1 numpy vector containing the samples from distributio
5     :param Y: An n x 1 numpy vector containing the samples from distributio
6     :param kernel_fn: supply the kernel function to use.
7     :return: the maximum mean discrepancy:
8     MMD(X,Y) = Expected value of k(X,X) + Expected value of k(Y,Y) - Expecte
9     where k is a kernel function
10    """
11    mmd = torch.mean(kernel_fn(X, X)) + torch.mean(kernel_fn(Y, Y)) - torch.
12    return mmd
13
14
15 def kernel(A, B):
16     """
17     A gaussian kernel on two arrays.
18     :param A: An n x d numpy matrix containing the samples from distributio
19     :param B: An n x d numpy matrix containing the samples from distributio
20     :return K: An n x n numpy matrix k, in which k_{i,j} = e^{-||A_i - B_j||
21     """
22    sigma = 1
23    K = torch.exp(-(torch.abs(A-B)**2 / (2.0 * sigma**2)))
24    return K

```

```
1 def compute_kl(a, b):
2     t_start = perf_counter()
3     kl = entropy(a, b)
4     t_stop = perf_counter()
5     return np.mean(kl), t_stop - t_start
6 def compute_emd(a, b):
7     kde_a = sklearn.neighbors.KernelDensity(kernel="gaussian", bandwidth=0.5).
8     a_samples = kde_a.score_samples(a)
9     kde_b = sklearn.neighbors.KernelDensity(kernel="gaussian", bandwidth=0.5).
10    b_samples = kde_b.score_samples(b)
11    dist = distance_matrix(np.expand_dims(a_samples, 1), np.expand_dims(b_samp
12    t_start = perf_counter()
13    emd_out = emd(np.exp(a_samples), np.exp(b_samples), dist)
14    t_stop = perf_counter()
15    return emd_out, t_stop - t_start
16 def compute_mmd(a, b):
17     t_start = perf_counter()
18     mmd_out = mmd(a, b, kernel)
19     t_stop = perf_counter()
20     return mmd_out, t_stop - t_start
21 def get_all(a, b):
22     kl1, kl1_time = compute_kl(a, b)
23     kl2, kl2_time = compute_kl(b, a)
24
25     emd_out, emd_time = compute_emd(a, b)
26
27     mmd_out, mmd_time = compute_mmd(a, b)
28
29     print('KL a, b: {} \t time: {}'.format(kl1, kl1_time))
30     print('KL b, a: {} \t time: {}'.format(kl2, kl2_time))
31     print('EMD: {} \t time: {}'.format(emd_out, emd_time))
32     print('MMD: {} \t time: {}'.format(mmd_out, mmd_time))
```

```
1 # PART 1
2 a = torch.rand((1000, 3))
3 b = torch.empty(1000, 3).normal_(mean=0,std=1)
4
5 #entropy(a, b), entropy(b,a) both will return [inf, inf, inf]
6
7 b[b==0] = 1e-8
8 #entropy(a, b), entropy(b,a) both still both return [inf, inf, inf]
9 b = torch.abs(b)
10 get_all(a, b)

KL a, b: 0.6029940843582153      time: 0.004756203999988884
KL b, a: 0.5953423976898193      time: 0.0003583929999990687
EMD: 947.8049488395836      time: 0.09825794299999646
MMD: 0.3636596202850342      time: 0.00039149700000962184
```

Question 5.1.1. Based on the above measures alone, which divergence seems most accurate?

I think either KL or MMD. EMD seems to be very high. However, since I know the values are uniform vs random, I imagine on average the distance between each point is < 0.5 , so I think MMD is the most accurate.


```

1 # PART 2
2 indices = torch.randperm(len(mnist_test))[:2000]
3
4 subset1 = torch.utils.data.Subset(mnist_test, indices[:1000])
5 subset2 = torch.utils.data.Subset(mnist_test, indices[1000:])
6
7 sub1 = []
8 for x, y in subset1:
9     sub1.append(x)
10
11 a = torch.cat(sub1, dim = 0).view(-1, 784)
12
13 sub2 = []
14 for x, y in subset2:
15     sub2.append(x)
16
17 b = torch.cat(sub2, dim = 0).view(-1, 784)
18
19 a[a==0] = 1e-8
20 b[b==0] = 1e-8
21
22 get_all(a, b)

```

KL a, b: 10.58070182800293	time: 0.028992849000019305
KL b, a: 10.555280685424805	time: 0.028606792000005044
EMD: 3.4907284985548466e-83	time: 0.09398919599999545
MMD: 0.11188805103302002	time: 0.010181770999963646

```

1 #PART 3
2 with open('p5_output.p', 'rb') as f:
3     b = pickle.load(f)
4 b = b.view(-1, 784)
5 # a[a==0] = 1e-8
6 b[b==0] = 1e-8
7 b = torch.abs(b)
8 get_all(a, b)

```

KL a, b: 2.190413236618042	time: 0.03119231400000899
KL b, a: 8.476593971252441	time: 0.03095107100000405
EMD: 2.5011979219961465e-75	time: 0.10797177799997826
MMD: 0.30129003524780273	time: 0.011146196999959557

Question 5.3.1. Which divergence or distance showed the greatest discrepancy between the comparison between real MNIST data and the comparison with the GAN?

In KL divergence when the GAN generated tensor is the first argument created the greatest discrepancy.

Question 5.3.2. Which of these information measures would you recommend for judging a GAN's output? Why?

I think using MMD would be the best. The discrepancy between the order of the KL divergence is too great. EMD seems to measure the distance to be incredibly small, which I know that cannot be the case. This leaves MMD which shows there is some discrepancy, but not too incredibly much, which seems about accurate for a GAN.

Question 5.3.3. How do the runtimes of these measures compare?

Based on the previous experiments, KL calculation seems to take much longer. EMD seems to be consistently the slowest and mmd was faster than KL in this case, but it doesn't seem to be the same for every experiment.